



Solving a system via QR Decomposition On NVIDIA GPUs provided Armadillo's sparse matrices (sp_mat)



Problem Description



- Create a highly sparse matrix with random values utilizing Armadillo's `sp_mat` class.
- Factorize and solve the matrix on a GPU utilizing `cusolver`.
- Compare the results with a sequential factorization.

Create a Highly Sparse Matrix utilizing Armadillo's 'sp_mat' Class

Easiest Part

- `int size = 512;`
- `sp_mat A = sprandn<sp_mat>(size, size, 0.03);`
- `vec b(size, fill::randu);`
 - For $Ax=b$

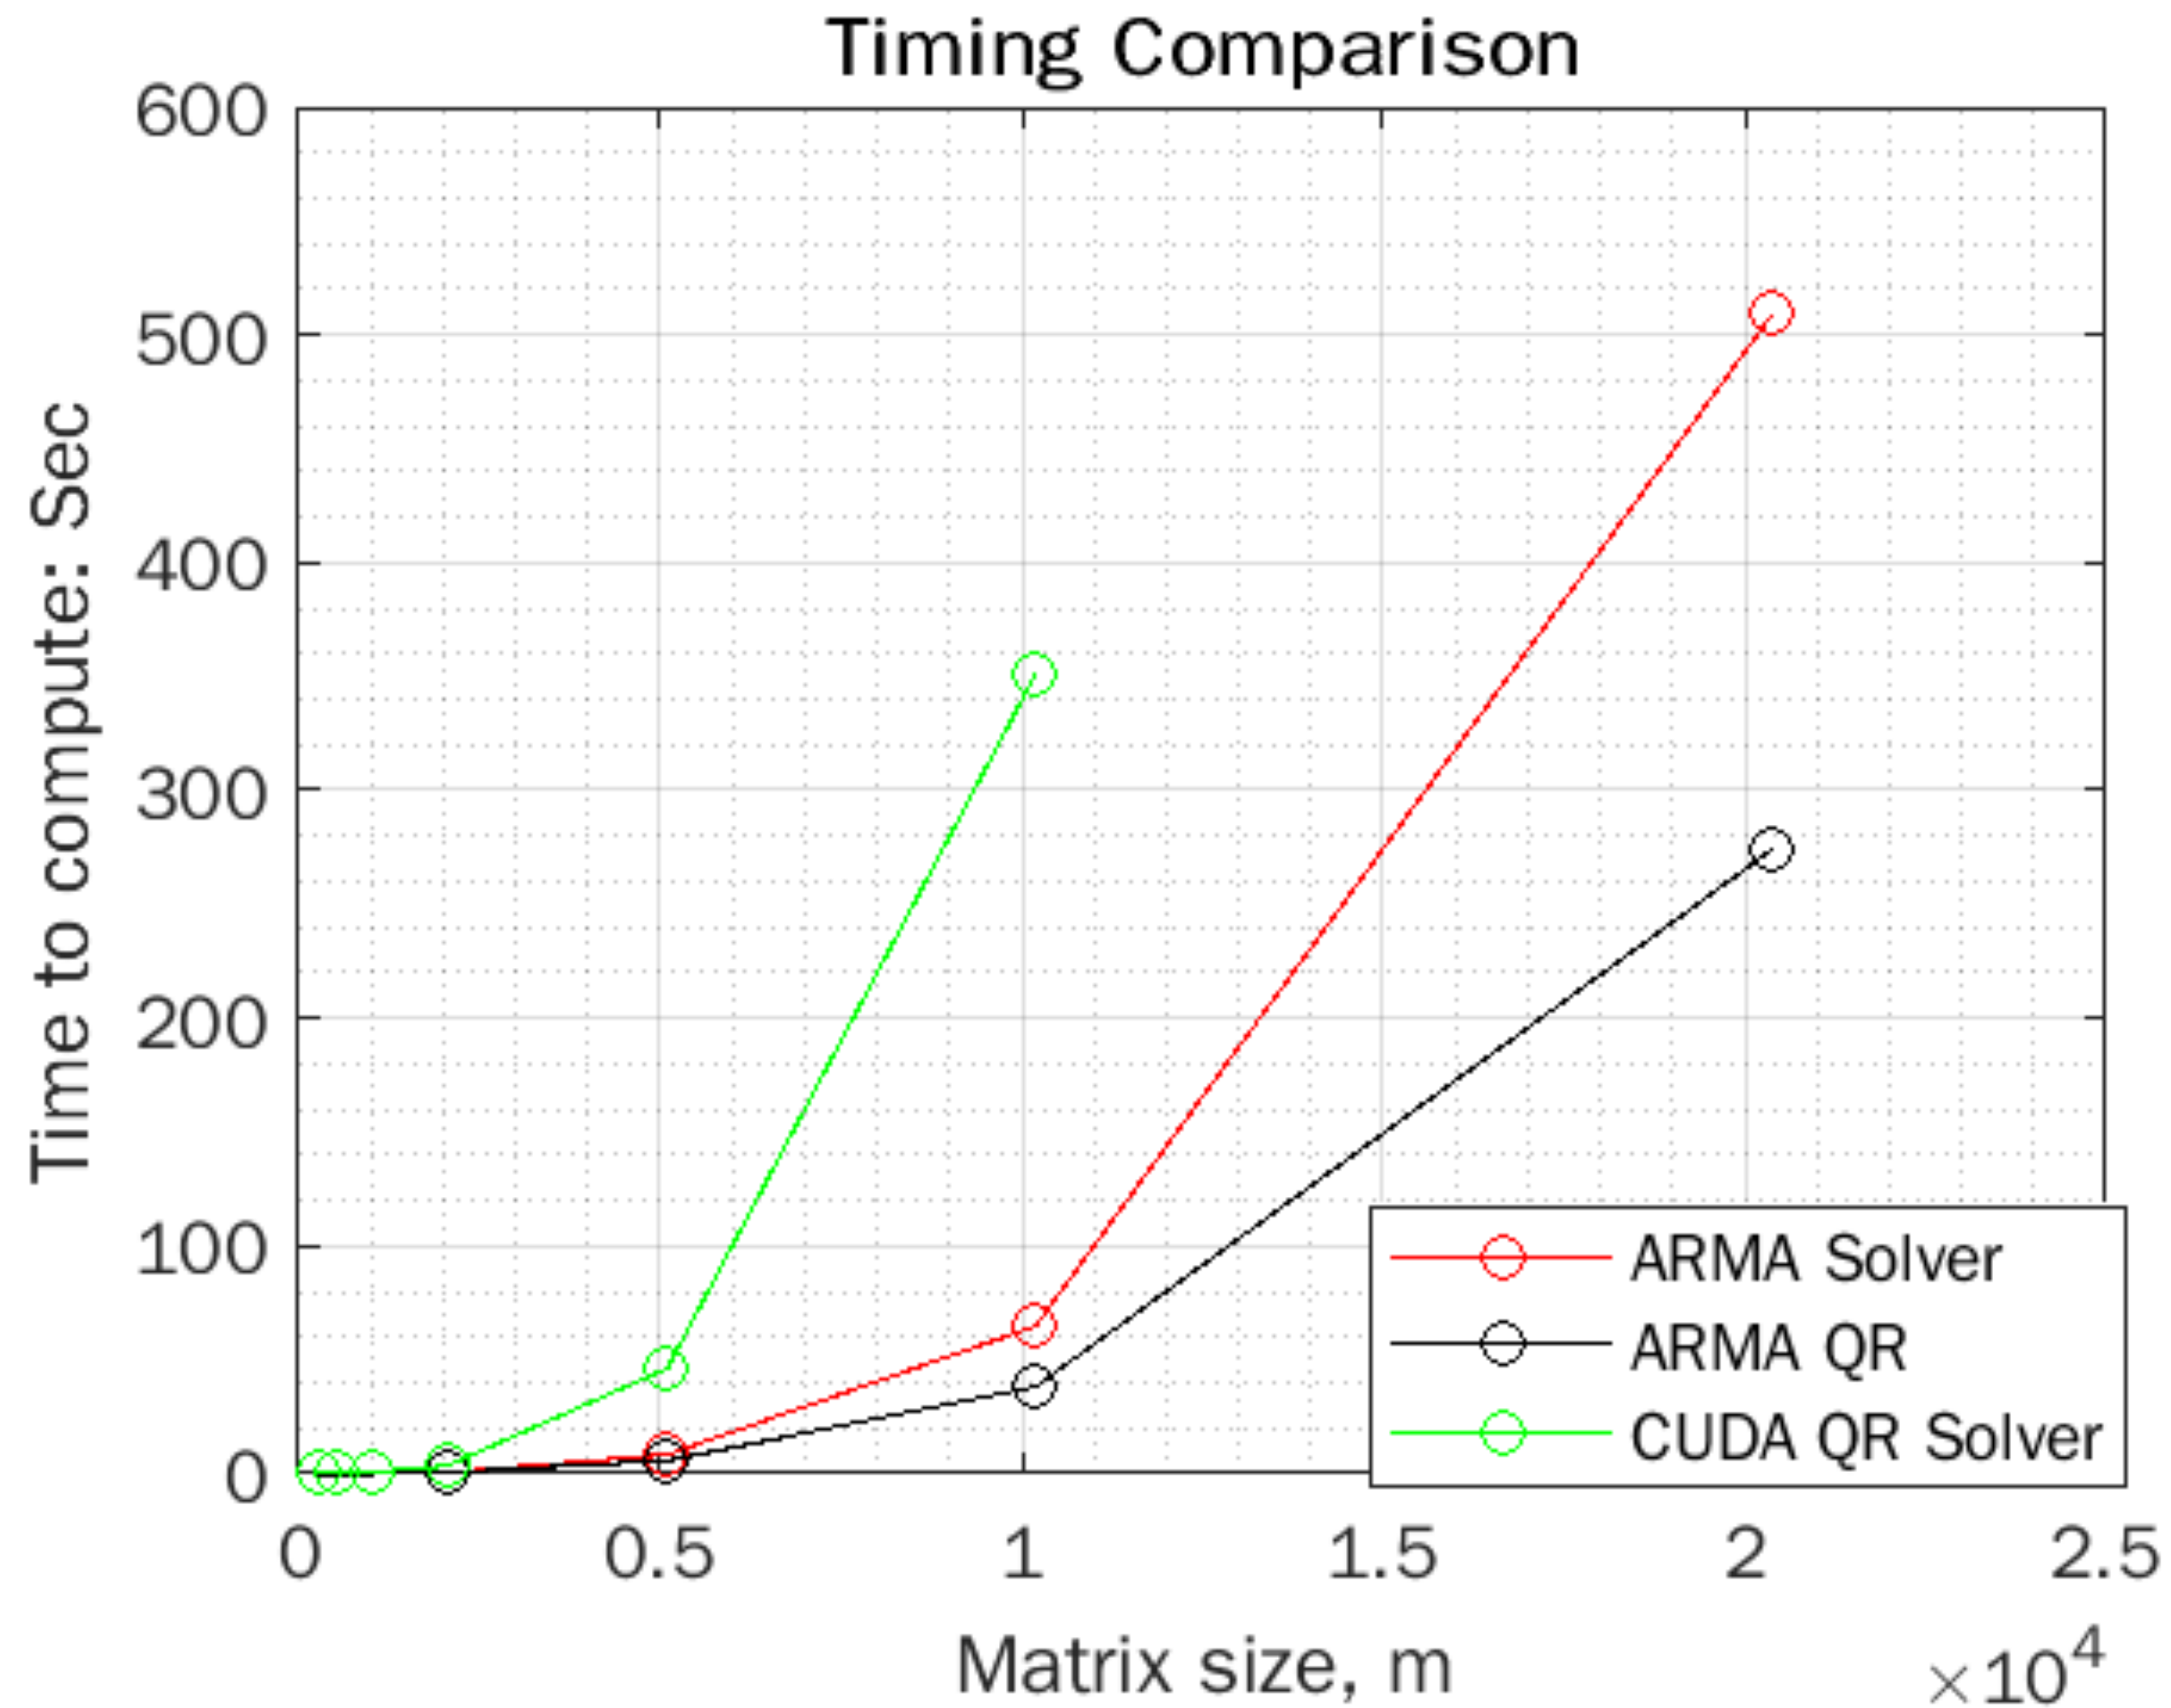


Factorize the matrix on a GPU utilizing cusolver.

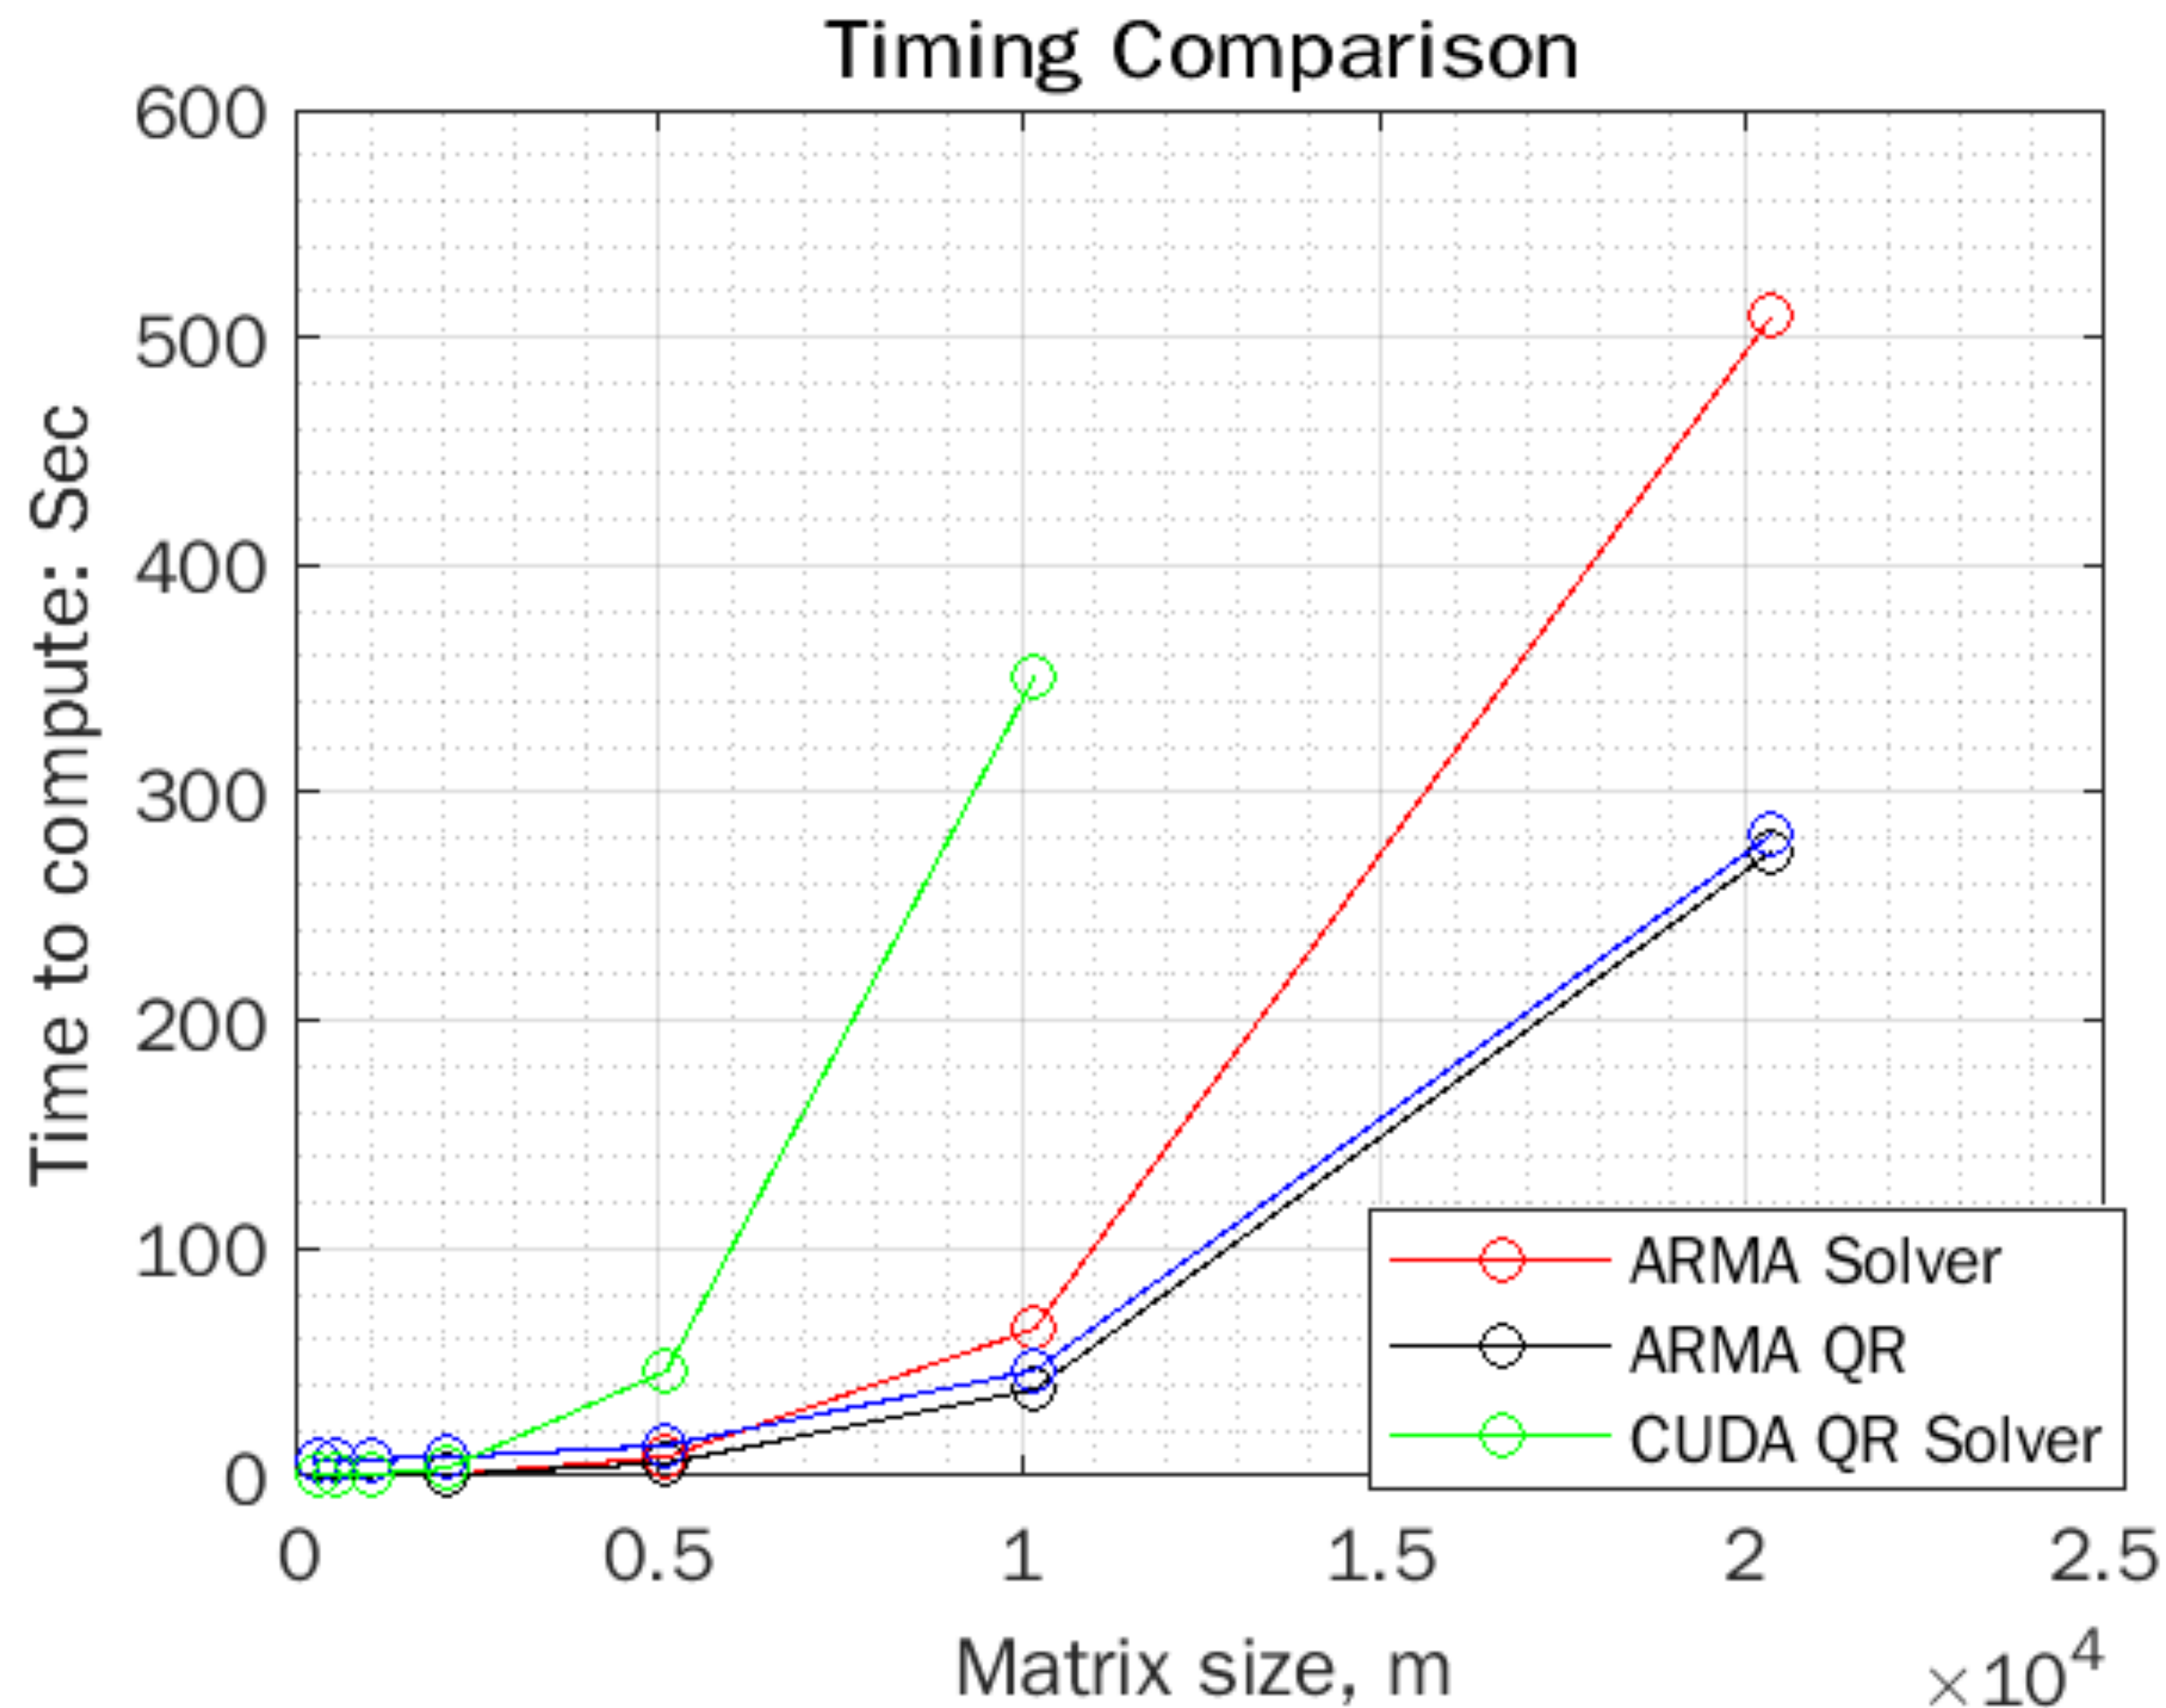
- Non-Trivial
 - Will go into greater depth



Compare the results with a sequential factorization.



Compare the results with a sequential factorization.



Factorize the matrix on a GPU utilizing cusolver.

Problems

- There is no public documentation for how Armadillo stores the `sp_mat` data besides that it is in CSC format.
- Armadillo stores its Sparse Matrices in CSC format while CUSOLVER requires CSR format.
- CUDA cusolverLU is only supported for host side computation and does not have GPU support.
 - Solving a sparse system in parallel is not a trivial problem.



Factorize the matrix on a GPU utilizing cusolver.

Problems

- It is not trivial to link and combine Armadillo Code and CUDA code
- Cusolver and cusparse libraries have poor documentation, and at least one inconsistency.



Sparse Matrix Storage

CSC vs CSR

- Because so many values in a sparse matrix are 0. We do not need to store them all.
- 2 common “compression” formats for sparse matrices



Sparse Matrix Storage

CSR

Ref[1]

- $V = [10\ 20\ 30\ 40\ 50\ 60\ 70\ 80]$
- $COL_INDEX = [0\ 1\ 1\ 3\ 2\ 3\ 4\ 5]$
- $ROW_INDEX = [0\ 2\ 4\ 7\ 8]$

$$\begin{pmatrix} 10 & 20 & 0 & 0 & 0 & 0 \\ 0 & 30 & 0 & 40 & 0 & 0 \\ 0 & 0 & 50 & 60 & 70 & 0 \\ 0 & 0 & 0 & 0 & 0 & 80 \end{pmatrix}$$



Sparse Matrix Storage

CSC

- $V = [10\ 20\ 30\ 40\ 50\ 60\ 70\ 80]$
- $COL_INDEX = [0\ 1\ 3\ 4\ 6\ 7\ 8]$
- $ROW_INDEX = [0\ 0\ 1\ 2\ 1\ 2\ 2\ 3]$

$$\begin{pmatrix} 10 & 20 & 0 & 0 & 0 & 0 \\ 0 & 30 & 0 & 40 & 0 & 0 \\ 0 & 0 & 50 & 60 & 70 & 0 \\ 0 & 0 & 0 & 0 & 0 & 80 \end{pmatrix}$$



Factorize the matrix on a GPU utilizing cusolver.

Problems

- There is no public documentation for how Armadillo stores the `sp_mat` data besides that it is in CSC format.
- Armadillo stores its Sparse Matrices in CSC format while CUSOLVER requires CSR format.
- CUDA cusolverLU is only supported for host side computation and does not have GPU support.
 - Solving a sparse system in parallel is not a trivial problem.



Factorize the matrix on a GPU utilizing cusolver.

Problems

- It is not trivial to link and combine Armadillo Code and CUDA code
- Cusolver and cusparse libraries have poor documentation, and at least one inconsistency.





Credit



- Special thanks to PHD Candidate Jared Brzenski for emotional support and details on problem scope.



References



- [1] [https://en.wikipedia.org/wiki/Sparse_matrix#Compressed_sparse_row_\(CSR,_CRS_or_Yale_format\)](https://en.wikipedia.org/wiki/Sparse_matrix#Compressed_sparse_row_(CSR,_CRS_or_Yale_format))